

SQL – Transactions

Carla Teixeira Lopes

Bases de Dados

Licenciatura em Engenharia Informática e Computação, FEUP+FCUP

Based on Jennifer Widom slides

Agenda

Introduction

Properties


Isolation levels

Motivation for transactions

Concurrent database access

Resilience to system failures

Concurrent Database Access



Attribute-level Inconsistency

Update College **Set** enr = enr + 1000 **Where** cName = 'Stanford'

concurrent with ...

Update College **Set** enr = enr + 1500 **Where** cName = 'Stanford'

	15000	

get; modify; put

$$15\ 000 + 2\ 500 = 17\ 500$$

$$15\ 000 + 1\ 000 = 16\ 000$$


$$15\ 000 + 1\ 500 = 16\ 500$$

Tuple-level Inconsistency

Update Apply Set major = 'CS' Where sID = 123

concurrent with ...

Update Apply Set dec = 'Y' Where sID = 123



sID	major	dec
123		

get; modify; put

both changes

One of the two changes

Table-level Inconsistency

Update Apply Set decision = 'Y'


Where sID In (Select sID From Student Where GPA > 3.9)

} S1

concurrent with ...

Update Student Set GPA = (1.1) * GPA Where sizeHS > 2500

} S2



Multi-statement Inconsistency


Insert Into Archive Select * From Apply Where decision = 'N';
Delete From Apply Where decision = 'N';

C1

concurrent with ...

Select Count(*) From Apply;
Select Count(*) From Archive;

C2



Concurrency goal

Execute sequence of SQL statements so they appear to be running in isolation

Simple solution: execute them in isolation


But want to enable concurrency whenever safe to do so

Multiprocessor system

Multithreaded system

Asynchronous I/O

Resilience to System Failures




Resilience to System Failures


Insert Into Archive

Select * From Apply Where decision = 'N';

Delete From Apply Where decision = 'N';




Resilience to System Failures



System-Failure Goal

Guarantee all-or-nothing execution, regardless of failures



Transactions

Solution for both concurrency and failures

A transaction is a sequence of one or more SQL operations treated as a unit

Transactions appear to run in isolation

If the system fails, each transaction's changes are reflected either entirely or not at all

Transactions: SQL standard

Transaction begins automatically on first SQL statement

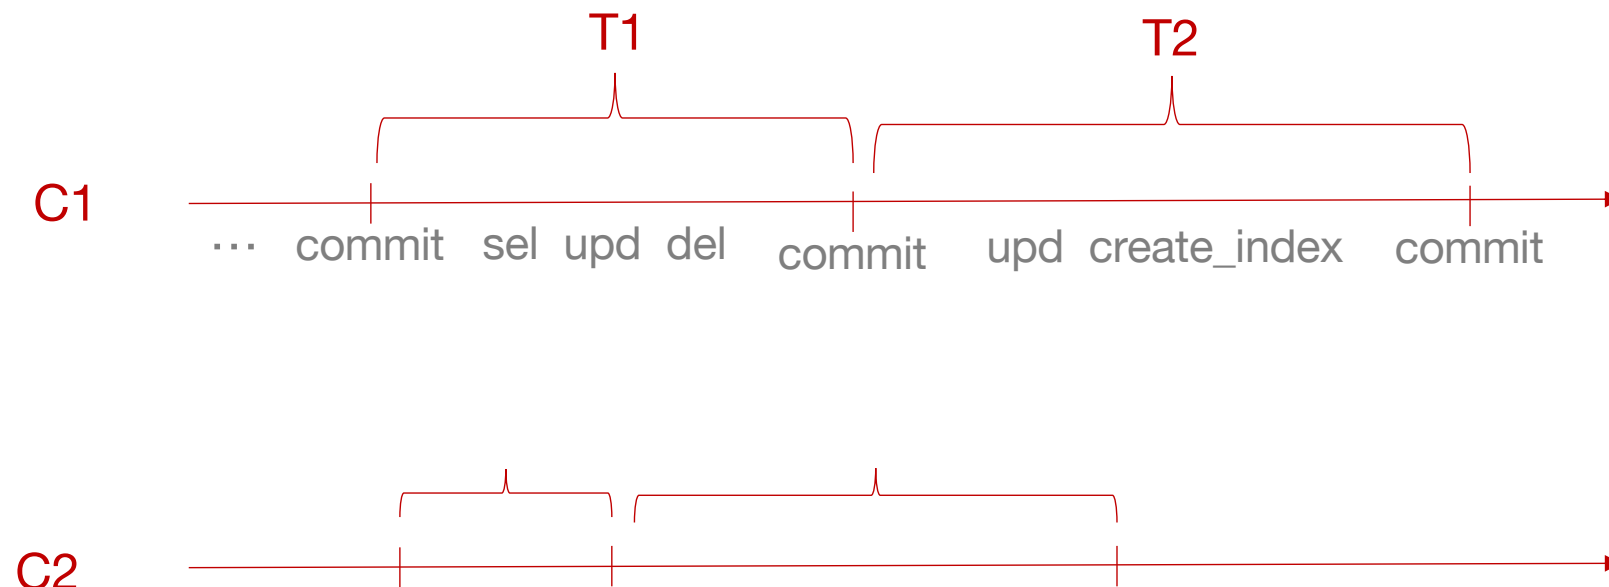
On “**commit**” transaction ends and new one begins

Current transaction ends on session termination

“**Autocommit**” turns each statement into transaction

Transactions

A transaction is a sequence of one or more SQL operations treated as a unit



Agenda

Introduction

Properties

Isolation levels

ACID Properties

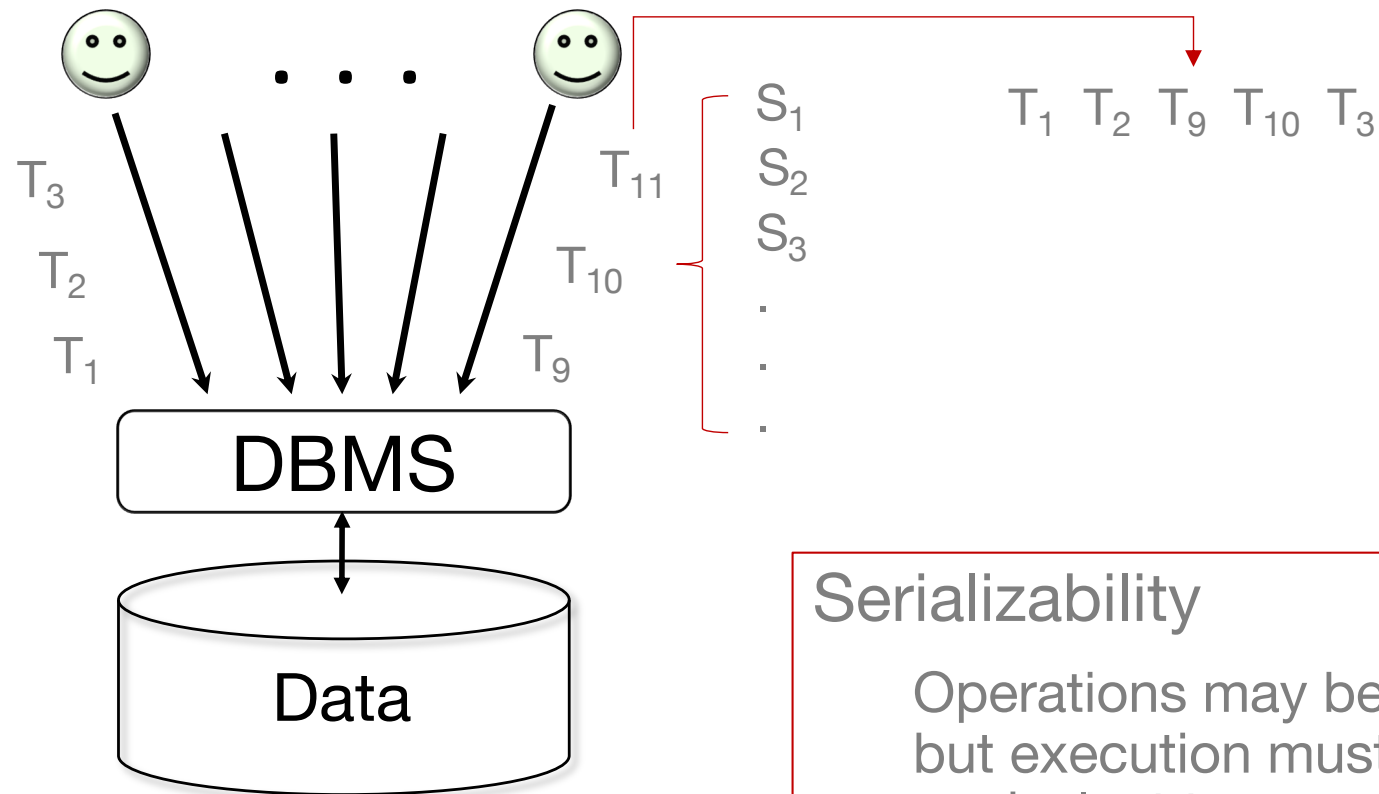
Atomicity 3

Consistency 4

Isolation 1

Durability 2

ACID Properties: Isolation



Serializability

Operations may be interleaved, but execution must be equivalent to some sequential (serial) order of all transactions

Attribute-level Inconsistency

Update College Set enr = enr + 1000 Where cName = 'Stanford' T_1

concurrent with ...

Update College Set enr = enr + 1500 Where cName = 'Stanford' T_2

If serializability is guaranteed

$$\begin{array}{l} T_1; T_2 \\ T_2; T_1 \end{array} \longrightarrow 15\ 000 \rightarrow 17\ 500$$

Tuple-level Inconsistency

Update Apply Set major = 'CS' Where sID = 123

T_1

concurrent with ...

Update Apply Set dec = 'Y' Where sID = 123

T_2

If serializability is guaranteed

$$\begin{array}{l} T_1; T_2 \\ T_2; T_1 \end{array} \longrightarrow \text{Both changes}$$

Table-level Inconsistency

Update Apply Set decision = 'Y'

Where sID In (Select sID From Student Where GPA > 3.9)

} T₁

concurrent with ...

Update Student Set GPA = (1.1) * GPA Where sizeHS > 2500

} T₂

If serializability is guaranteed

T₁; T₂
T₂; T₁

Order
matters

DBMS don't guarantee the
exact sequential order if the
transactions are being
issued at the same time

Multi-statement Inconsistency

Insert Into Archive Select * From Apply Where decision = 'N';
Delete From Apply Where decision = 'N';

T_1

concurrent with ...

Select Count(*) From Apply;
Select Count(*) From Archive;


T_2

If serializability is guaranteed

$T_1; T_2$
 $T_2; T_1$ \longrightarrow Order matters

ACID Properties: Durability


If system crashes after transaction commits, all effects of transaction remain in database



ACID Properties: Atomicity

Each transaction is “all-or-nothing”, never left half done

Using a logging mechanism, partial effects of transactions at the time of crash are undone



Transaction Rollback (= Abort)

Undoes partial effects of transaction

Can be system- or client-initiated

```
Begin Transaction;  
<get input from user>  
SQL commands based on input  
<confirm results with user>  
If ans='ok' Then Commit; Else Rollback;
```


Transactions should
be constructed to run
quickly

→ Not wait arbitrary
amounts of time

Locking

Only undoes effects on the data itself

ACID Properties: Consistency



Each client, each transaction:

Can assume all constraints hold when transaction begins

Must guarantee all constraints hold when transaction ends

Serializability

Constraints always hold




Agenda

Introduction

Properties

Isolation levels

ACID Properties: Isolation




Serializability

Operations may be interleaved, but execution must be equivalent to some sequential (serial) order of all transactions

Disadvantages

Overhead in locking
Reduction in concurrency

ACID Properties: Isolation



Weaker “Isolation Levels”

Read Uncommitted

Read Committed

Repeatable Read

Serializable

Weak
↓
Strong

↓ Overhead in locking

↑ Concurrency


↓ Consistency Guarantees

Isolation Levels

Per transaction

It does not affect the behaviour of any other transaction

Specific to Reads



Dirty Reads

“Dirty” data item: written by an uncommitted transaction

Update College Set enr = enr + 1000 Where cName = 'Stanford'

T₁

concurrent with ...

Select avg(enr) From College

T₂



If read before T1 commits, this value is known as dirty

Assume there is a commit at the end of each box

Dirty Reads – Example 2

Update Student Set GPA = (1.1) * GPA Where sizeHS > 2500

T₁

concurrent with ...

Select GPA From Student Where sID=123

T₂

concurrent with ...

Update Student Set sizeHS=2600 Where sID=234

T₃

Where can we have dirty data items?

There are no
dirty reads
within the same
transaction

Read Uncommitted

A transaction may perform dirty reads

Update Student Set GPA = (1.1) * GPA Where sizeHS > 2500

T₁

concurrent with ...

Select avg(GPA) From Student

T₂

If transactions are serializable

T1; T2 or

T2; T1

Read Uncommitted

Update Student Set GPA = (1.1) * GPA Where sizeHS > 2500

T₁

concurrent with ...

Set Transaction Isolation Level Read Uncommitted;
Select avg(GPA) From Student;

T₂

We don't have serializable behaviour

We might don't care that much about consistency

Read Committed

A transaction may **not** perform dirty reads

Still does not guarantee global serializability

Update Student Set GPA = (1.1) * GPA Where sizeHS > 2500

T₁

concurrent with ...

Set Transaction Isolation Level Read Committed;

Select avg(GPA) From Student

Select max(GPA) From Student

T₂

Repeatable Read

A transaction may **not** perform dirty reads

An item read multiple times cannot change value

Still does not guarantee global serializability

Update Student Set GPA = (1.1) * GPA Where sizeHS > 2500;

Update Student Set sizeHS=1500 Where sID = 123;

T₁

concurrent with ...

Set Transaction Isolation Level Repeatable Read;

Select avg(GPA) From Student

Select avg(sizeHS) From Student

T₂

Repeatable Read

A transaction may **not** perform dirty reads

An item read multiple times cannot change value

But a relation *can* change: “phantom” tuples

Insert into Student [100 new tuples]

T_1

Phantom tuples

concurrent with ...

Set Transaction Isolation Level Repeatable Read;

Select avg(GPA) From Student

Select max(GPA) From Student

T_2

Repeatable Read

A transaction may **not** perform dirty reads

An item read multiple times cannot change value

But a relation *can* change: “phantom” tuples

Delete from Student [100 new tuples]

T₁

concurrent with ...

Set Transaction Isolation Level Repeatable Read;

Select avg(GPA) From Student

Select max(GPA) From Student

T₂

Once read, values get locked and deletion is not possible in the middle of T₂

Read Only Transactions

Helps system optimize performance

Independent of isolation level

Not going to perform modifications to the database within the transaction



Set Transaction Read Only;

Set Transaction Isolation Level Repeatable Read;

Select avg(GPA) From Student

Select max(GPA) From Student

weak		dirty reads	nonrepeatable reads	phantoms
	Read Uncommitted	Y	Y	Y
	Read Committed	N	Y	Y
	Repeatable Read	N	N	Y
	Serializable	N	N	N
strong				

Isolation Levels: Summary

Standard default: Serializable

Weaker isolation levels

Increased concurrency + decreased overhead = increased performance

Weaker consistency guarantees

Some systems have default Repeatable Read

Isolation level per transaction

Each transaction's reads must conform to its isolation level

Kahoot time!

Any doubts?

Readings

Jeffrey Ullman, Jennifer Widom, A first course in
Database Systems 3rd Edition

Section 6.6 – Transactions in SQL